# Programming Camp Day 2

Augustus Kmetz

2025-09-11

# Welcome back!
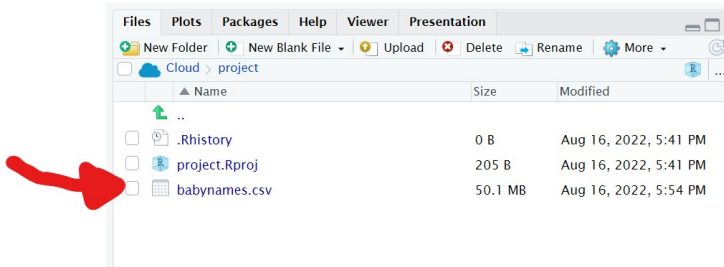
Plans for today:

- Importing data
- Transforming data
- Tidy data

# babynames.csv

babynames.csv

- names and sex of babies born in the US from 1880 to 2017
- data compiled by the Social Security Administration
- 1.9M rows

# Number of children with the name Augustus



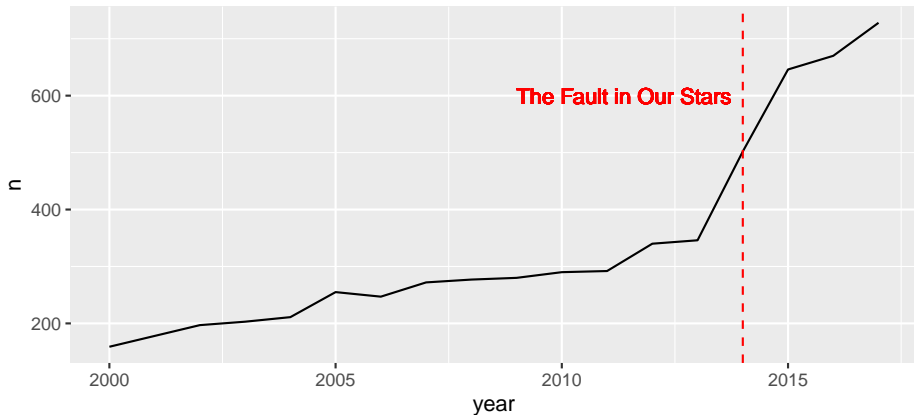**Q**: Why was there a spike in children with the name Augustus?

# Number of children with the name Augustus



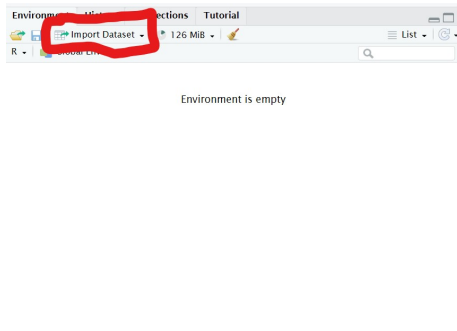**Q**: Why was there a spike in children with the name Augustus?

**Answer:** The movie "The Fault in Our Stars" was released in 2014.

# Plausible Explanation: Release of the movie "The Fault in Our Stars"

# Import Data

# Import a dataset (the wrong way)

# Import a dataset (the wrong way)
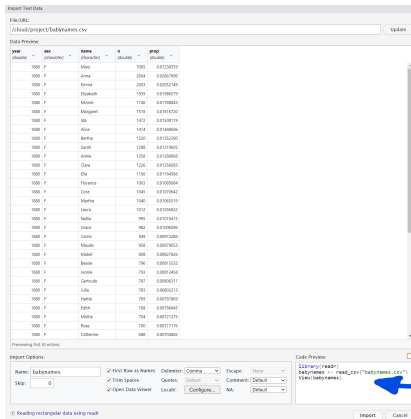


**Q:** But is this reproducible?

# Working directory

- The current working directory is the location which R is currently pointing to
- Whenever you try to read or save a file without specifying the path explicitly that is where R will "look"
- To see the current working directory use getwd()
- To change the working directory use setwd("path/to/directory")
- When you open a R file that folder automatically becomes the wd
- **Best practice** is to specify files paths relative to the main folder of a project

# Paths and directory names

- R inherits its file and folder naming conventions from UNIX and uses forward slashes for the directories, e.g. `/data/raw/FEMA/`
- Backslashes serve a different purpose - they are used as escape characters and to isolate special characters
- To avoid problems, directory names should **NOT** contain spaces and special characters

# Back to babynames

- A common text file format is a comma delimited text file: .csv
- These files use comma as a column separator

```
year,sex,name,n,prop
1880,F,Mary,7065,0.0724
1880,F,Anna,2604,0.0267
```

- To read these files use the following command

```
my_data <- read_csv("path/to/filename.csv")
```

# The readr package

- `readr` is for reading rectangular text data into R

- Very fast and easy to customize how file is parsed

- `readr` supports several file formats with general `read_<...>()` functions:
  - `read_csv()`: comma-separated files
  - `read_tsv()`: tab-separated files
  - `read_delim()`: general delimited files
  - `read_fwf()`: fixed-width files
  - `read_table()`: tabular files where columns are separated by white-space
  - `read_log()`: web log files

# Importing other types of data

Rectangular data:

- package `haven` reads SPSS, Stata and SAS files
- package `readxl` reads excel files (both .xls and .xlsx)

Hierarchical data:

- `jsonlite` for json (common for browser-server communications)
- `xml2` for XML (common for textual data in web services)

And many many more...

# Back to babynames

```
babynames <- read_csv("data/babynames.csv")
```

```
babynames
```

```
# A tibble: 1,924,665 x 5
    year sex   name           n    prop
   <dbl> <chr> <chr>      <int>   <dbl>
 1  1880 F     Mary        7065  0.0724
 2  1880 F     Anna        2604  0.0267
 3  1880 F     Emma        2003  0.0205
 4  1880 F     Elizabeth   1939  0.0199
 5  1880 F     Minnie      1746  0.0179
 6  1880 F     Margaret    1578  0.0162
 7  1880 F     Ida         1472  0.0151
 8  1880 F     Alice       1414  0.0145
 9  1880 F     Bertha      1320  0.0135
10  1880 F     Sarah       1288  0.0132
# i 1,924,655 more rows
```

# Back to babynames

Number of Boys Named Augustus over Time
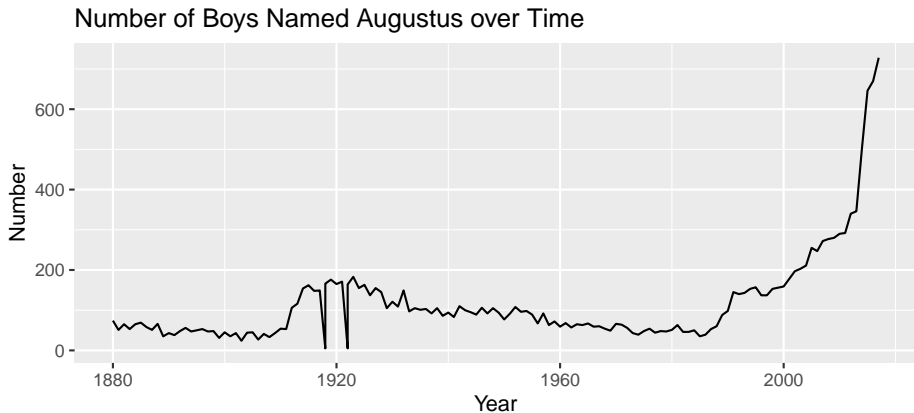


**Q:** Which geom?

# Plotting babynames

What happens if we plot the data?

```
ggplot(babynames) +
  geom_line(mapping = aes(x = year, y = prop))
```

# Plotting babynames

Try a different geom?

```
ggplot(babynames) +
  geom_point(mapping = aes(x = year, y = prop))
```

# How to isolate?

babynames

| year | sex | name | n | prop |
|------|-----|------|-----|------|
| 1880 | M | John | 9655 | 0.0815 |
| 1880 | M | William | 9532 | 0.0805 |
| 1880 | M | James | 5927 | 0.0501 |
| 1880 | M | Charles | 5348 | 0.0451 |
| 1880 | M | Garrett | 13 | 0.0001 |
| 1881 | M | John | 8769 | 0.081 |

| year | sex | name | n | prop |
|------|-----|------|-----|------|
| 1880 | M | Garrett | 13 | 0.0001 |
| 1881 | M | Garrett | 7 | 0.0001 |
| … | … | Garrett | … | … |

# Transforming data

# Transforming data

The `dplyr` package is part of the `tidyverse` which:

- introduces a **grammar** for transforming tabular data
- is fast on data frames (written in C++) speed of C with ease of R
- intuitive to write and easy to read esp when using chaining syntax

# dplyr verbs (functions)

dplyr verbs handle the vast majority of your data needs:

- `filter()` - for picking observations by their values
- `select()` - for picking variables by their names
- `arrange()` - for reordering rows
- `mutate()` - for creating new variables with functions on existing variables
- `summarise()` - for collapsing many values down to a single summary

All of the above can be done using ugly and slow base R functions

# The structure of `dplyr` functions

All verbs work similarly:

- The first argument is a tibble (or data frame)
- The subsequent ones describe what to do, using the variable names
- The results is a new tibble

# select()

Extract columns by name

```
select(.data, name_col_1, name_col_2, ...)
```

```
select(babynames, name, prop)
```

babynames

| year | sex | name | n | prop |
|------|-----|---------|------|--------|
| 1880 | M | John | 9655 | 0.0815 |
| 1880 | M | William | 9532 | 0.0805 |
| 1880 | M | James | 5927 | 0.0501 |
| 1880 | M | Charles | 5348 | 0.0451 |
| 1880 | M | Garrett | 13 | 0.0001 |
| 1881 | M | John | 8769 | 0.081 |

➡

| name | prop |
|---------|--------|
| John | 0.0815 |
| William | 0.0805 |
| James | 0.0501 |
| Charles | 0.0451 |
| Garrett | 0.0001 |
| John | 0.081 |

# select() helpers

**:** - select range of columns

```
select(babynames, year:name)
```

**-** - select every column but …

```
select(babynames, -prop)
```

**starts_with()** - select columns that start with…

```
select(babynames, starts_with("n"))
```

**ends_with()** - select every column that ends with …

```
select(babynames, ends_with("e"))
```

And many more!

# Quiz

Which of these is NOT a way to select the **name** and **n** columns together?

```
select(babynames, -c(year, sex, prop))
```

```
select(babynames, starts_with("n"))
```

```
select(babynames, ends_with("n"))
```

# filter()

Extract rows that meet logical criteria

```
filter(.data, condition_1, condition_2, ...)
```

```
filter(babynames, name == "Garrett")
```

# Logical tests

| | |
|---|---|
| x < y | Less than |
| x > y | Greater than |
| x == y | Equal to |
| x <= y | Less than or equal to |
| x >= y | Greater than or equal to |
| x != y | Not equal to |
| x %in% y | Group membership |

# Logical tests

```
x <- 1
x >= 2
```

```
[1] FALSE
```

```
x <- c(1,2,3)
x >= 2
```

```
[1] FALSE  TRUE  TRUE
```

# Missing Values

NA stands for a missing value

Handling missing/infinite/divide by 0 values can be a bit tricky.

What is the result?

```
1 == 1
1 == NA
NA == NA
1 + Inf
1 + NaN
Inf == Inf
```

# Logical tests

| | |
|---|---|
| x  < y | Less than |
| x  > y | Greater than |
| x  == y | Equal to |
| x  <= y | Less than or equal to |
| x  >= y | Greater than or equal to |
| x  != y | Not equal to |
| x  %in% y | Group membership |
| is.na(x) | Is NA |
| !is.na(x) | Is not NA |

Same syntax works for `Inf` and `NaN`: `is.inf()`, `is.nan()`

# Your turn!

Use `filter`, `babynames` and the logical operators to:

- Make a plot of your name's popularity over time (or linguistically similar name)
- Which name had the highest ever number/proportion of babies?

# Two common mistakes

1. Using = instead of ==

```
filter(babynames, name = "Augustus")
filter(babynames, name =="Augustus")
```

2. Forgetting quotes

```
filter(babynames, name == Augustus)
filter(babynames, name =="Augustus")
```

# Boolean operators

| | |
|---|---|
| *a* & *b* | and |
| *a* \| *b* | or |
| xor(*a*, *b*) | exactly or |
| !*a* | not |
| ( ) | To group tests . & evaluates before \| |

## Your turn!

Use Boolean operators to alter the code below to return only the rows that contain:

- Boys named Leslie
- Names that were used by exactly 5 or 6 children in 1880
- Names that are one of Anakin, Leia, Luke

```
filter(babyname, name == "Sea" | name == "Anemone")
```

# Two more common mistakes

3. Collapsing multiple tests into one

```
filter(babynames, 10 < n < 20)
filter(babynames, 10 < n, n < 20)
```

4. Stringing together many tests (when you could use %in%)

```
filter(babynames, n == 5 | n == 6 | n == 7 )
filter(babynames, n %in% c(5, 6, 7))
```

# arrange()

Order rows from smallest to largest values

```
arrange(.data, order_first, order_second, ...)
```

```
arrange(babynames, n)
```

order rows from largest to smallest values

```
arrange(babynames, desc(n))
```

**Q**: INSERT QUESTION HERE

# Steps

Goal: create a data set called `top_5_M_2015` that contains the 5 most popular boy baby names from 2015

Steps

1. filter to boys born in 2015
2. select the name and n columns
3. arrange these columns so that the most popular names appear near the top
4. use `head` to take the top 5

# Steps

```r
top_5_M_2015 <- filter(babynames, year == 2015, sex == "M")
top_5_M_2015 <- select(top_5_M_2015, name, n)
top_5_M_2015 <- arrange(top_5_M_2015, desc(n))
top_5_M_2015 <- head(top_5_M_2015, n = 5)
```

```r
top_5_M_2015 <-
  head(arrange(select(filter(babynames, year == 2015, sex == '
  name, n), desc(n)), n = 5)
```

# The pipe operator %>%

```
filter(babynames, year == 2015, sex == "M")
babynames %>% filter(year == 2015, sex == "M")
```

Passes results on left into first argument of function on right

In pipe notation you use:

- x %>% f(y) rather than f(x, y)
- x %>% f(y) %>% h(z) rather than h(f(x,y),z)

This is similar to Unix pipes

# Pipes

```
top_5_M_2015 <-
  babynames %>%
  filter(year == 2015, sex == "M") %>%
  select(name, n) %>%
  arrange(desc(n)) %>%
  head(n = 5)
```

Try to use %>% to write a sequence of functions that make a plot of your name's popularity over time (or linguistically similar name)

# Pipes

```
babynames %>%
  filter(name == "Augustus", sex == "M") %>%
  ggplot() +
  geom_line(aes(x = year, y = n))
```

# What are the most popular names?

Do we have tools to:

1. Calculate the total number of children with each name?
2. Calculate the number of unique names over time?
3. Calculate the share of babies captured by the top 10 names?

...not quite yet

## summarise()

compute table of summaries

```
summarise(.data,
          summary_stat_1 = function_1(variable_1),
          summary_stat_2 = function_2(variable_2))
```

```
babynames %>%
  summarise(total = sum(n), max = max(n))
```

```
babynames %>%
  filter(name == "Augustus") %>%
  summarise(total = sum(n), max = max(n))
```

**Q**: How can we do this for each name?

# summarise()

compute table of summaries

```
summarise(.data,
          summary_stat_1 = function_1(variable_1),
          summary_stat_2 = function_2(variable_2))
```

```
babynames %>%
  summarise(total = sum(n), max = max(n))
```

```
babynames %>%
  filter(name == "Augustus") %>%
  summarise(total = sum(n), max = max(n))
```

**Q**: How can we do this for each name?

**Answer:** Using the group_by() function (coming up next).

# group_by()

Groups cases by common values

```
group_by(.data, first_variable, second_variable, ...)
```

```
babynames %>%
  group_by(sex) %>%
  summarise(total = sum(n))
```

ungroup() - removes grouping criteria from a data frame

```
babynames %>%
  group_by(sex) %>%
  ungroup() %>%
  summarise(total = sum(n))
```

# Elementary but useful summary functions

Aggregation functions return one value per group

- `min(x), median(x), max(x), quantile(x, p)`
- `n(), n_distinct(), sum(x), mean(x)`
- `sum(x > 10), mean(x > 0)`
- `sd(x), var(x)`

Window functions return multiple values per group

- `top_n()`
- `lead(), lag()`

# Your turn!

Partner up

1. For each year calculate the number of distinct names by sex. Plot these time series.
2. Plot the share of babies with a name among the top 10 names over time by sex

# Distinct names Over time



Number of Distinct Baby Names over Time

Note: Only names with at least five records are included in the SSA data.

# Distinct names Over time

```
babynames %>%
  group_by(year, sex) %>%
  summarise(n_names = n_distinct(name)) %>%
  ggplot(aes(x = year, y = n_names, color = sex)) +
  geom_line() +
  labs(x = "Year",
       y = "Number",
       title = "Number of Distinct Baby Names over Time",
       caption = "Note: Only names with at least five records
```

# Top share over time



Share of Babies with Top 10 Names over Time

# Top share over time

```r
babynames %>%
  group_by(year, sex) %>%
  arrange(prop) %>%
  top_n(10, wt = prop) %>%
  summarise(total_prop = sum(prop)) %>%
  ggplot(aes(x = year, y = total_prop, color = sex)) +
  geom_line() +
  labs(x = "Year",
       y = "Share",
       title = "Share of Babies with Top 10 Names over Time")
```

## mutate()

Create new columns

```
mutate(.data,
        new_var_1 = function_1(variable_1),
        new_var_2 = function_2(variable_2))
```

```
babynames %>%
  mutate(percent = round(prop * 100, 2))
```

Any vectorized function can be used with mutate() including:

- arithmetic operators (+, -, /, *, etc)
- logical operators (<, <=, >=, ==, !=)
- logarithmic and exponential transformations (log, log10, exp)
- offsets (lead, lag)
- cumulative rolling aggregates (cumsum, cumprod, cummin, cummax)
- ranking (min_rank, percent_rank)

# Other useful `dplyr` functions

- `rename()` - change the name of variables
- `count()` - shortcut for `group_by() %>% summarise(n = n())`
- `distinct()` - 1 row for each variable value
- `sample_n()`, `sample_frac()` - random sample of tibble
- `pull()` - extract a single column as a vector
- `slice()` - subset rows using their position

# Names that "flash in the pan"

What names experience a dramatic rise/fall in popularity?

# Your turn

Step 1: find the names that experience a collapse in popularity

1. Filter to names that have at least 1% babies with that name (for a sex)
2. For each name calculate the percent change in `prop` over time
3. Take the top 10 most dramatic collapses over time

Step 2: plot the popularity over time for those names

# Names that "flash in the pan"



Names that Flash in the Pan

# Names that "flash in the pan"

```
babynames_flash <-
  babynames %>%
  group_by(name, sex) %>%
  mutate(max_prop = max(prop)) %>%
  filter(max_prop > 0.01) %>%
  mutate(change_in_pop = (prop/lag(prop, 1) - 1)) %>%
  summarise(min_change = min(change_in_pop, na.rm = T)) %>%
  group_by(sex) %>%
  arrange(min_change) %>%
  slice(1:5) %>%
  pull(name)
```

# Names that "switch" - optional homework

Challenge: Come up with a way to identify names that "switch"

# dplyr recap

 Extract variables with **select()**

 Extract cases with **filter()**

 Arrange cases, with **arrange()**.

 Make tables of summaries with **summarise()**.

 Make new variables, with **mutate()**.